

HW8

Jingshi

4/2/2017

7.9 #2

a)

When λ (the smoothing factor) is extremely large and $m = 0$, $g^{(0)}$ (which equals to g) is forced to converge to 0. Therefore, $\hat{g} = 0$.

b)

When λ (the smoothing factor) is extremely large and $m = 1$, $g^{(1)}$ (which is the first derivative of g) is forced to converge to 0. Therefore, $\hat{g} = c$ where c stands for a constant. For example, $\hat{g} = 9$.

c)


When λ (the smoothing factor) is extremely large and $m = 2$, $g^{(2)}$ (which is the second derivative of g) is forced to converge to 0. Therefore, $\hat{g} = a \times x + b$ where a and b are both stands for constants. For example, $\hat{g} = 9 \times x + 3$.

7.9 #6

b)

```
library(boot)
library(ISLR)
#head(Wage)

# perform cross-validation to choose the optimal number of cuts
set.seed(6)
# Divide into 2 to 15 cuts
cross.validation <- rep(NA, 15)
for (i in 2:15) {
  Wage$cut <- cut(Wage$age, i)
  fit <- glm(wage ~ cut, data = Wage)
  # 12 fold cross validation
  cross.validation[i] <- cv.glm(Wage, fit, K = 12)$delta[1]
}
plot(2:15, cross.validation[-1], xlab = "Number of Cuts", ylab = "Testing MSE", type = "o")
points(which.min(cross.validation), cross.validation[which.min(cross.validation)], col = "green",
       pch = 1, cex = 2)
```




As the above plot shows, when number of cuts is 11, the testing mean squared error is the lowest among the other cuts. Thus, the optimal number of cuts is 11.

```
# fit a step function
model.step <- glm(wage ~ cut(age, 11), data = Wage)
# take a look at the summary
summary(model.step)
```

```
##
## Call:
## glm(formula = wage ~ cut(age, 11), data = Wage)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -99.89  -24.77   -5.43   15.46  203.22
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      70.580      3.340  21.133 < 2e-16 ***
## cut(age, 11)(23.6,29.3]  23.719      4.048   5.860 5.14e-09 ***
## cut(age, 11)(29.3,34.9]  35.768      3.921   9.122 < 2e-16 ***
## cut(age, 11)(34.9,40.5]  47.001      3.785  12.418 < 2e-16 ***
## cut(age, 11)(40.5,46.2]  49.399      3.752  13.167 < 2e-16 ***
## cut(age, 11)(46.2,51.8]  46.571      3.833  12.150 < 2e-16 ***
## cut(age, 11)(51.8,57.5]  46.781      3.936  11.887 < 2e-16 ***
## cut(age, 11)(57.5,63.1]  50.142      4.279  11.719 < 2e-16 ***
## cut(age, 11)(63.1,68.7]  37.296      6.946   5.369 8.50e-08 ***
## cut(age, 11)(68.7,74.4]  26.414      8.020   3.294  0.001 **
## cut(age, 11)(74.4,80.1]  20.491     13.064   1.569  0.117
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 1595.04)
```

```
##
## Null deviance: 5222086 on 2999 degrees of freedom
## Residual deviance: 4767575 on 2989 degrees of freedom
## AIC: 30651
##
## Number of Fisher Scoring iterations: 2
# preperation for the plot
seq.age <- seq(from = range(Wage$age)[1], to = range(Wage$age)[2])
prediction <- predict(model.step, list(age = seq.age))

# make a plot of the fit obtained.
plot(Wage$wage ~ Wage$age)
lines(seq.age, prediction, col = "green", lwd = 3)
```



7.9 #9

a)


```
library(MASS)
# use the poly() function to fit a cubic polynomial regression to predict nox using dis.
model.poly <- lm(nox ~ poly(dis, 3), data = Boston)
# report the regression output.
summary(model.poly)
```

```
##
## Call:
## lm(formula = nox ~ poly(dis, 3), data = Boston)
##
## Residuals:
```

##	Min	1Q	Median	3Q	Max
----	-----	----	--------	----	-----

```
## -0.121130 -0.040619 -0.009738 0.023385 0.194904
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.554695   0.002759 201.021 < 2e-16 ***
## poly(dis, 3)1 -2.003096   0.062071 -32.271 < 2e-16 ***
## poly(dis, 3)2  0.856330   0.062071  13.796 < 2e-16 ***
## poly(dis, 3)3 -0.318049   0.062071  -5.124 4.27e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.06207 on 502 degrees of freedom
## Multiple R-squared:  0.7148, Adjusted R-squared:  0.7131
## F-statistic: 419.3 on 3 and 502 DF, p-value: < 2.2e-16
```

```
# plot the resulting data.
plot(Boston$nox ~ Boston$dis)
# Plot the polynomial fits.
seq.dis <- seq(from = range(Boston$dis)[1], to = range(Boston$dis)[2], by = 0.01)
prediction <- predict(model.poly, data.frame(dis=seq.dis))
lines(seq.dis, prediction, col = "green", lwd = 3)
```



b)

```
library(data.table)
# create a vector with 10 elements in order to record residual sum of squares
# for each one of the 10 polynomial degrees.
myRss <- rep(NA, 10)
# fit 10 polynomial models and record their associated residual sum of squares.


for (i in 1:10) {
```

```


# fit a polynomial model.
model.poly <- lm(nox ~ poly(dis, i), data = Boston)
# plot the polynomial fit
plot(Boston$nox ~ Boston$dis, main=paste("Polynomial Degree: ", i))
seq.dis <- seq(from = range(Boston$dis)[1], to = range(Boston$dis)[2], by = 0.01)
prediction <- predict(model.poly, data.frame(dis=seq.dis))
lines(seq.dis, prediction, col = "green", lwd = 3)
# record the associated residual sum of squares.
myRss[i] <- sum(model.poly$residuals^2)
}

```


Polynomial Degree: 1




Polynomial Degree: 2




Polynomial Degree: 3




Polynomial Degree: 4




Polynomial Degree: 5




Polynomial Degree: 6



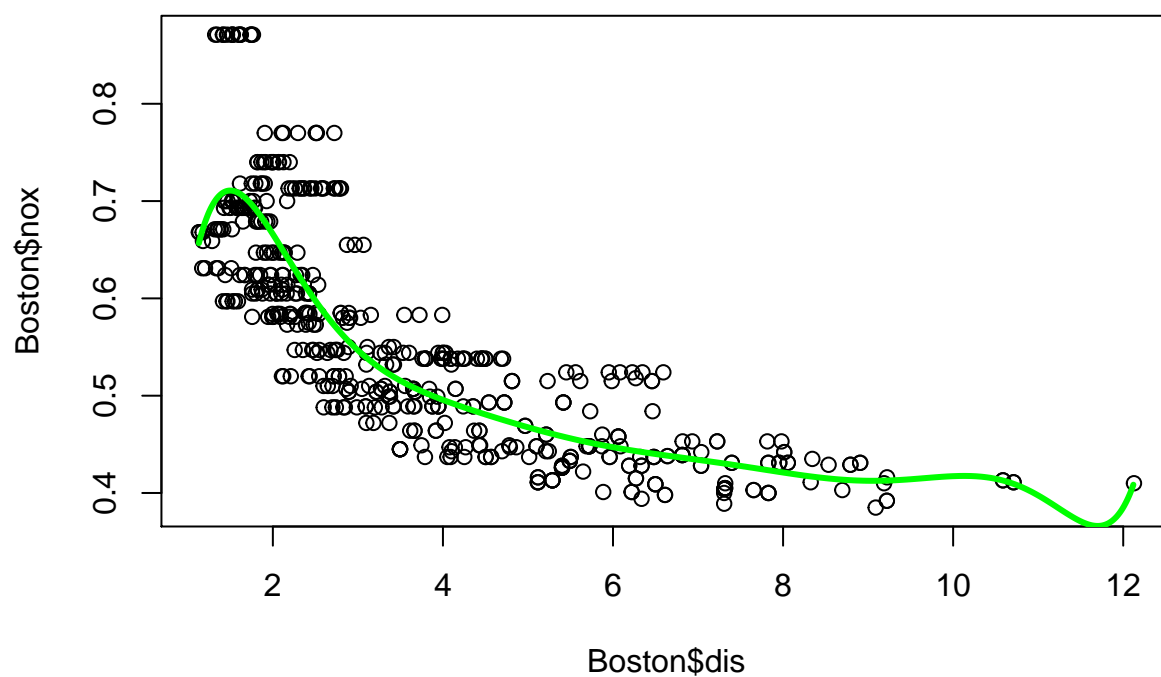
Polynomial Degree: 7




Polynomial Degree: 8



Polynomial Degree: 9




Polynomial Degree: 10



report the RSS's by plotting them out.

```
plot(1:10, myRss, xlab = "Polynomial Degrees", ylab = "Residual Sum of Squares", type = "o")
```



```
RSS<-myRss
```

create a table

```
myTable<-data.table(RSS)
```

take a look at the table


```
myTable[, Polynomial_Degree := .GRP, by = list(c(1,2,3,4,5,6,7,8,9,10))]
```

```
##          RSS Polynomial_Degree
## 1: 2.768563          1
## 2: 2.035262          2
## 3: 1.934107          3
## 4: 1.932981          4
## 5: 1.915290          5
## 6: 1.878257          6
## 7: 1.849484          7
## 8: 1.835630          8
## 9: 1.833331          9
## 10: 1.832171         10
```

According to the above table and the plots of RSS show, as the polynomial degree increase, RSS decrease. When polynomial degree is 10, the RSS is the smallest.

c)

```
set.seed(6)
library(boot)
# perform cross-validation for polynomial degrees from 1 to 10.
testMSE <- rep(NA, 10)
for (i in 1:10) {
  model.poly <- glm(nox ~ poly(dis, i), data = Boston)
  testMSE[i] <- cv.glm(Boston, model.poly, K = 12)$delta[1]
}
plot(1:10, testMSE, xlab = "Polynomial Degree", ylab = "Test Mean Squared Error", type = "o")
points(which.min(testMSE), testMSE[which.min(testMSE)], col = "green",
       pch = 1, cex = 2)
```



As you can see from the above plot, the test mean squared error is minimized when polynomial degree is 4. Therefore, the optimal degree for the polynomial is 4.


7.9 #9

d)

```
library(splines)
# fit a regression spline
model.lm <- lm(nox ~ bs(dis, df=4), data = Boston)
#bs(Boston$dis, df=4)
# Report the output for the fit
summary(model.lm)

##
## Call:
## lm(formula = nox ~ bs(dis, df = 4), data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.124622 -0.039259 -0.008514  0.020850  0.193891
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.73447    0.01460   50.306 < 2e-16 ***
## bs(dis, df = 4)1 -0.05810    0.02186   -2.658  0.00812 **
## bs(dis, df = 4)2 -0.46356    0.02366  -19.596 < 2e-16 ***
## bs(dis, df = 4)3 -0.19979    0.04311   -4.634  4.58e-06 ***
## bs(dis, df = 4)4 -0.38881    0.04551   -8.544 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.06195 on 501 degrees of freedom
## Multiple R-squared:  0.7164, Adjusted R-squared:  0.7142
## F-statistic: 316.5 on 4 and 501 DF,  p-value: < 2.2e-16

# preparation for the below plot
prediction <- predict(model.lm, data.frame(dis = seq.dis))
# Plot the resulting fit.
plot(Boston$nox ~ Boston$dis)
lines(seq.dis, prediction, col = "green", lwd = 3)
```




I choose the knots based on $df = 4$. Thus the knots are 1.296, 3.20745 and 12.1265 (this is obtained from `bs(Boston$dis, df=4)` and the output is too long to show here), each term in the fitted model are significant since their p-values are much smaller than $\alpha=0.05$. Moreover, the model's adjusted R-squared is high at 0.71.


e)

```
myRss <- rep(NA, 20)
for (i in 3:20) {
  model.lm <- lm(nox ~ bs(dis, df = i), data = Boston)
  # plot the fit
  plot(Boston$nox ~ Boston$dis, main=paste("Degree of Freedom: ", i))
  seq.dis <- seq(from = range(Boston$dis)[1], to = range(Boston$dis)[2], by = 0.01)
  prediction <- predict(model.lm, data.frame(dis=seq.dis))
  lines(seq.dis, prediction, col = "green", lwd = 3)
  myRss[i] <- sum(model.lm$residuals^2)
}
```


Degree of Freedom: 3




Degree of Freedom: 4




Degree of Freedom: 5




Boston\$dis
Degree of Freedom: 6




Degree of Freedom: 7




Boston\$dis
Degree of Freedom: 8




Degree of Freedom: 9




Degree of Freedom: 10




Degree of Freedom: 11




Degree of Freedom: 12




Degree of Freedom: 13




Degree of Freedom: 14




Degree of Freedom: 15




Degree of Freedom: 16




Degree of Freedom: 17




Boston\$dis
Degree of Freedom: 18




Degree of Freedom: 19



Degree of Freedom: 20



```
plot(3:20, myRss[-c(1, 2)], xlab = "Degrees of freedom", ylab = "Residual Sum Squares", type = "o")
```



```
RSS<-myRss
# create a table
myTable<-data.table(RSS)
# take a look at the table
myTable[, Degree_of_Freedom := .GRP, by = list(c(1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17,18,19,20))]
```

```
##      RSS Degree_of_Freedom
##  1:    NA                1
##  2:    NA                2
##  3: 1.934107              3
##  4: 1.922775              4
##  5: 1.840173              5
##  6: 1.833966              6
##  7: 1.829884              7
##  8: 1.816995              8
##  9: 1.825653              9
## 10: 1.792535             10
## 11: 1.796992             11
## 12: 1.788999             12
## 13: 1.782350             13
## 14: 1.781838             14
## 15: 1.782798             15
## 16: 1.783546             16
## 17: 1.779789             17
## 18: 1.775838             18
## 19: 1.774487             19
## 20: 1.776727             20
```

As the output shows, the RSS decrease generally as degrees of freedom increase. Although there are some fluctuations (for instance, from 10 degrees of freedom to 11 degrees of freedom, there is a slightly increase in RSS), the general trend of RSS value is decreasing.


7.9 #10

a)

```
library(leaps)

## Warning: package 'leaps' was built under R version 3.3.2

set.seed(2441)
attach(College)
# Record the total number of predictors.
n=length(College[,])-1
# Use a 70-30 split into training and test data.
data.train <- sample(length(Outstate), length(Outstate) *0.7)
mytest <- College[-data.train, ]
mytrain <- College[data.train, ]
# perform forward stepwise selection on the training set.
model <- regsubsets(Outstate ~ ., data = mytrain, nvmax = n, method = "forward")
# record the summary of the model.
mysummary <- summary(model)
par(mfrow = c(1, 3))
# Plot Adjusted R-squared vs # of variables.
plot(mysummary$adjr2, xlab = "# of variables", ylab = "Adjusted R-squared", type = "o")
# plot the maximum point.
points(which.max(mysummary$adjr2), mysummary$adjr2[which.max(mysummary$adjr2)], col = "green",
       pch = 1, cex = 2)
# +/- 0.1 standard deviations of the maximum point.
abline(h = max(mysummary$adjr2)+ 0.1 * sd(mysummary$adjr2), lty = 3, col = "purple")
abline(h = max(mysummary$adjr2) - 0.1 * sd(mysummary$adjr2), lty = 3, col = "purple")
# Plot BIC vs # of variables.
plot(mysummary$bic, xlab = "# of variables", ylab = "BIC", type='o')
# plot the minimum point.
points(which.min(mysummary$bic), mysummary$bic[which.min(mysummary$bic)], col = "green",
       pch = 1, cex = 2)
# +/- 0.1 standard deviations of the minimum point.
abline(h = min(mysummary$bic) + 0.1 * sd(mysummary$bic), lty = 3, col = "purple")
abline(h = min(mysummary$bic) - 0.1 * sd(mysummary$bic), lty = 3, col = "purple")
# Plot Cp vs # of variables.
plot(mysummary$cp, xlab = "# of variables", ylab = "C_p", type = "o")
# plot the minimum point.
points(which.min(mysummary$cp), mysummary$cp[which.min(mysummary$cp)], col = "green",
       pch = 1, cex = 2)
# +/- 0.1 standard deviations of the minimum point.
abline(h = min(mysummary$cp) + 0.1 *sd(mysummary$cp), lty = 3, col = "purple")
abline(h = min(mysummary$cp) - 0.1 *sd(mysummary$cp), lty = 3, col = "purple")
```

As the output shows, adjusted R-squared achieved its highest at 15 variables, BIC achieved its lowest at 7 variables, Cp achieved its lowest at 14 variables. These are optimal number of variables for each of the criterion. However, they are different. Thus, I choose 11 variables model as the satisfactory model that uses just a subset of the predictors because this model's Adjusted R-squared, BIC and Cp are within 0.1 standard deviation of their optimal values of each of the three criterion.

Take a look at the predictors that are selected from 11 variables model.

```
names(coef(model, id = 11))
```

```
## [1] "(Intercept)" "PrivateYes" "Apps" "Accept" "Top25perc"
## [6] "Room.Board" "Personal" "PhD" "S.F.Ratio" "perc.alumni"
## [11] "Expend" "Grad.Rate"
```

b)

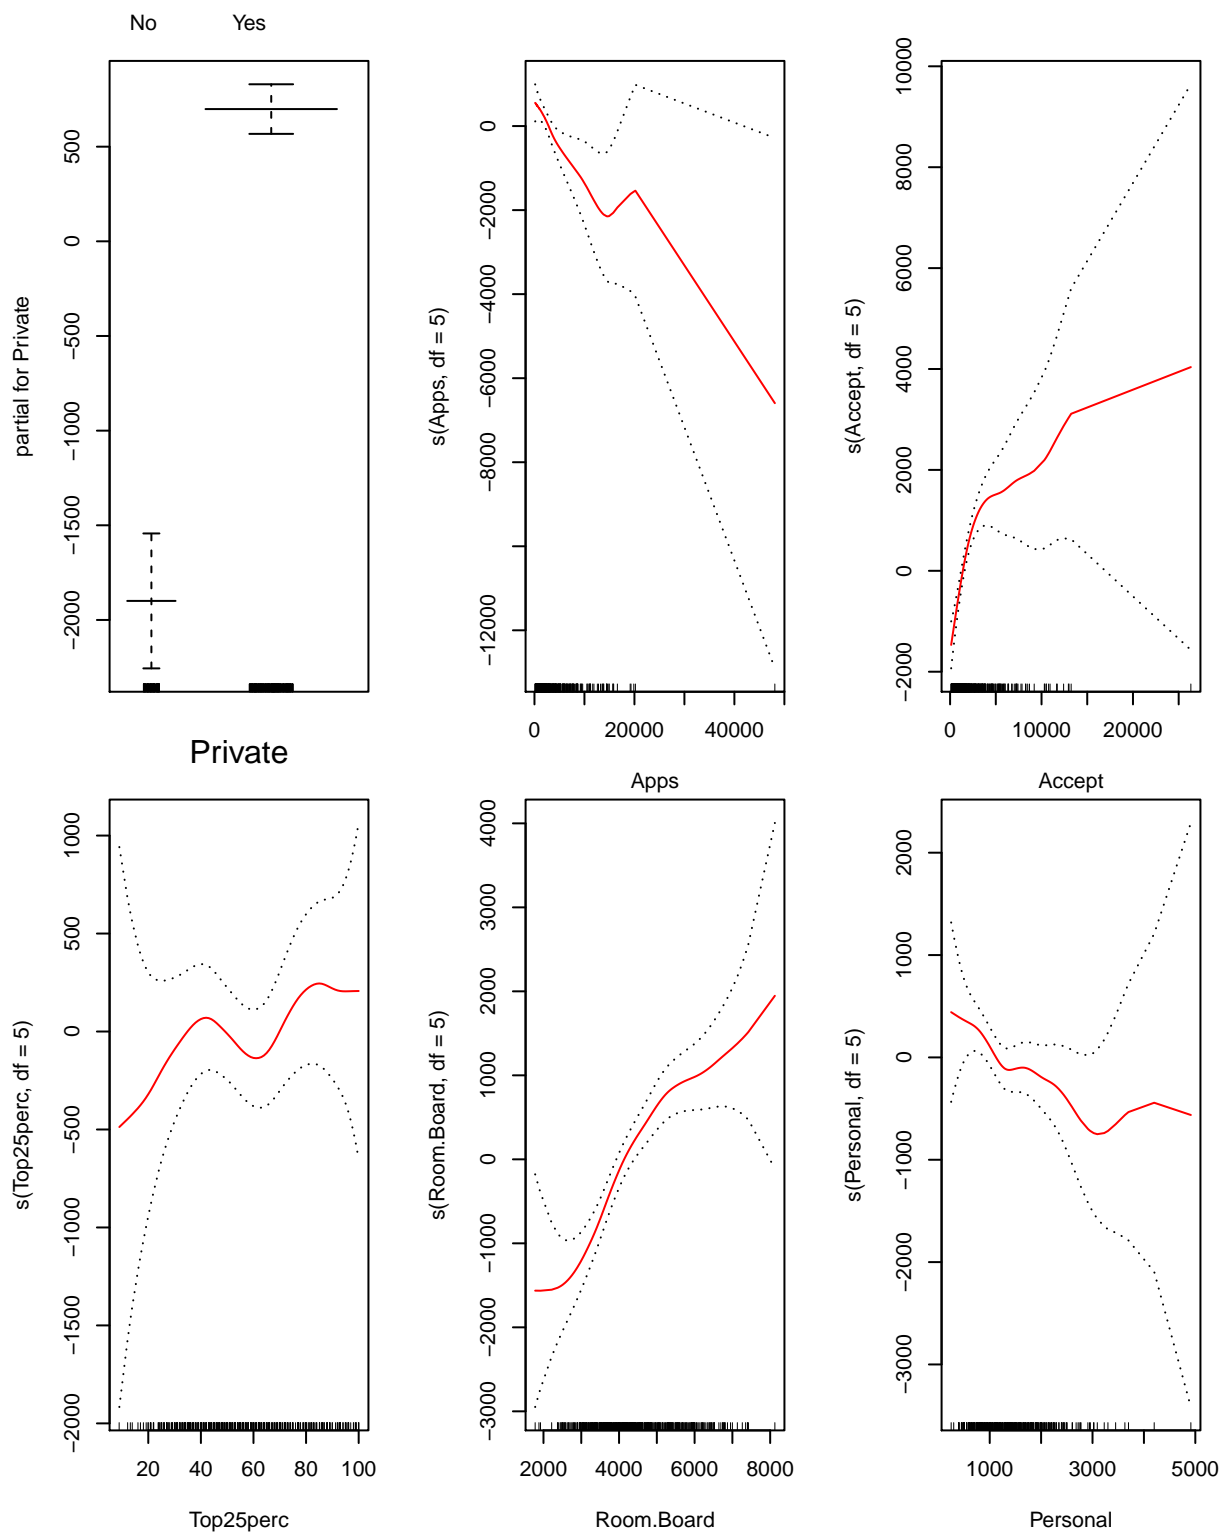
```
#install.packages("gam")
```

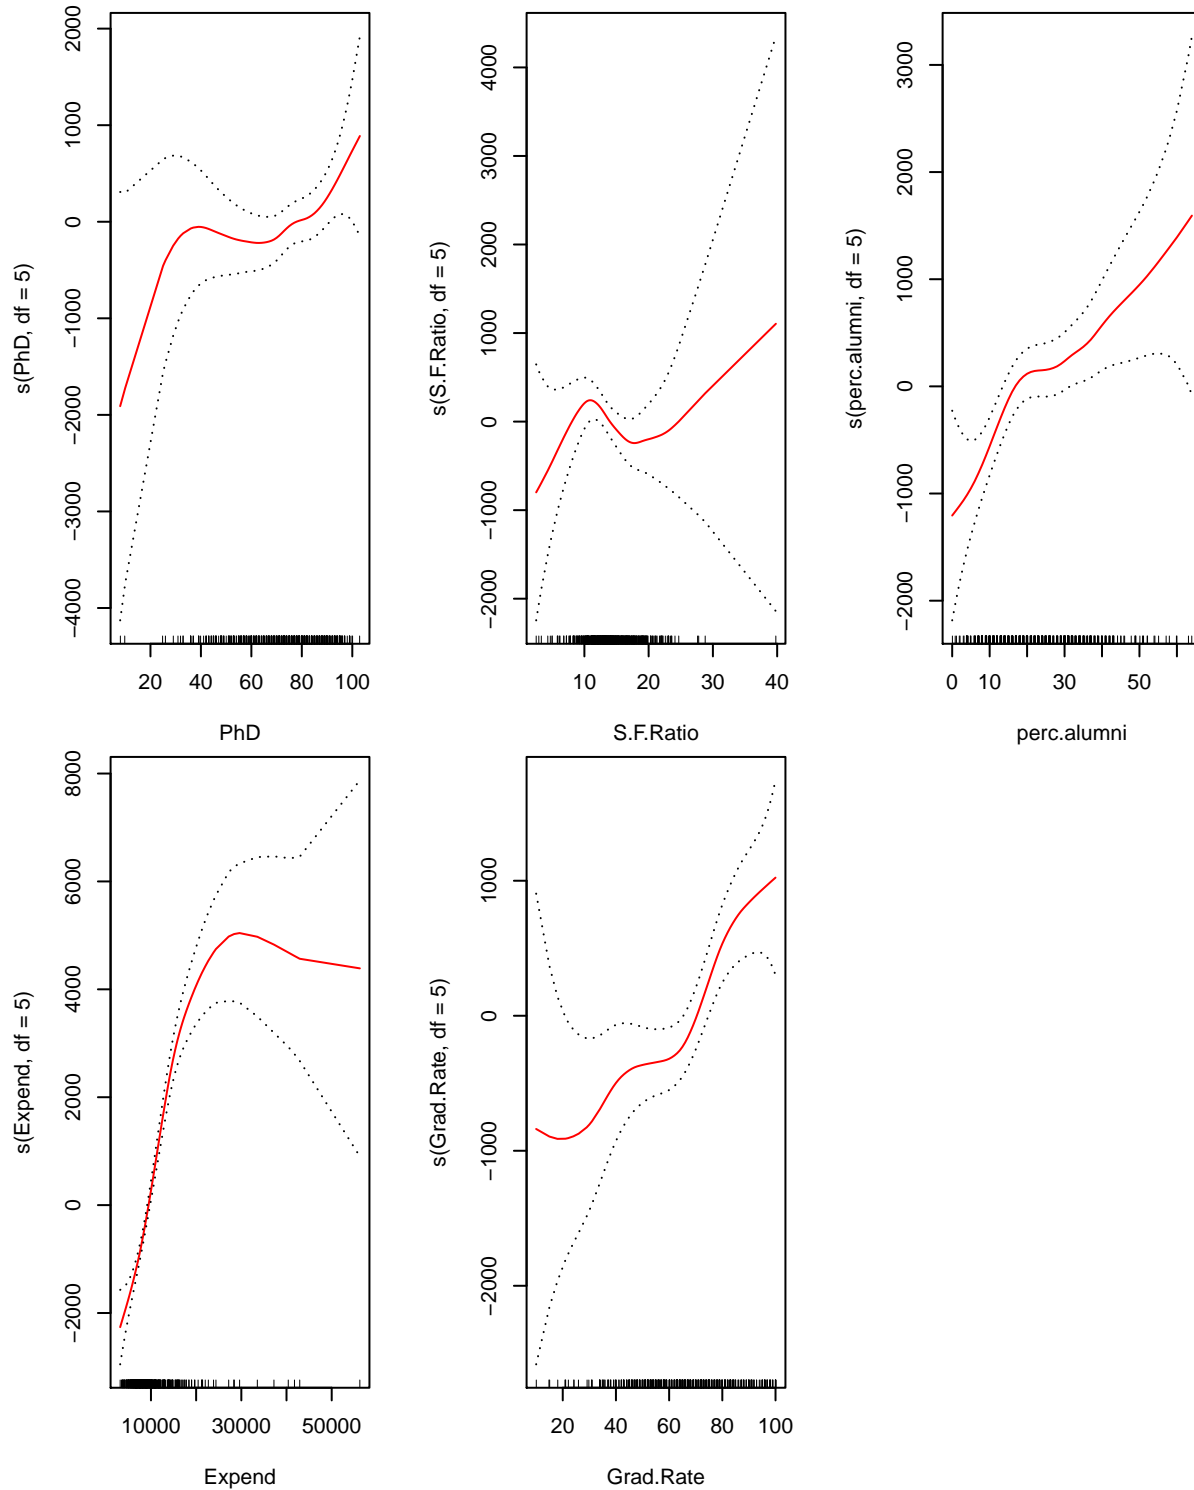
```
library(gam)
```

```
## Loading required package: foreach
```

```
## Loaded gam 1.14
```

```
model <- gam(Outstate ~ Private + s(Apps, df = 5) + s(Accept, df = 5) + s(Top25perc, df = 5) +
             s(Room.Board, df = 5) + s(Personal, df = 5) + s(PhD, df = 5) + s(S.F.Ratio, df=5)
             +s(perc.alumni, df=5)+s(Expend, df=5)+s(Grad.Rate, df=5), data=mytrain)
par(mfrow = c(1, 3))
plot(model, se = T, col = "red")
```





As the output shows, holding other variables constant, when 'Private' is yes, outstate tuition is high compared to when 'Private' is no. Holding other variables constant, when 'Apps' increase, outstate tuition decrease generally but fluctuate slightly in the middle. Holding other variables constant, when 'Accept' increase, outstate tuition increases in general. Holding other variables constant, when 'Top25perc' increase, outstate tuition increase moderately in general but fluctuate in the middle. Holding other variables constant, when 'Room Board' increase, outstate tuition increase in general. Holding other variables constant, when 'Personal' increase, outstate tuition decrease in general but fluctuate in the middle. Holding other variables constant,

when 'PhD' increase, outstate tuition increase in general but fluctuate in the middle. Holding other variables constant, when 'S.F.Ratio' increase, outstate tuition increase in general but fluctuate in the middle. Holding other variables constant, when 'perc alumni' increase, outstate tuition increase in general. Holding other variables constant, when 'Expend' increase, outstate tuition increase at first then followed by a decreasing. Holding other variables constant, when 'Grad Rate' increase, outstate tuition increase in general but fluctuate slightly in the middle.

c)

```
# Calculate root mean squared of 11 predictors GAM model.
prediction <- predict(model, mytest)
error<-mean((mytest$Outstate - prediction)^2)
rms<- sqrt(mean((mytest$Outstate - prediction)^2))
cat("Root Mean Squared of 11 predictors GAM model:", rms)

## Root Mean Squared of 11 predictors GAM model: 2070.563

# Calculate root mean squared of standard linear model with all predictors included.
model.lm<-lm(Outstate ~., data=mytrain)
prediction.lm <- predict(model.lm, mytest)
rms.lm<- sqrt(mean((mytest$Outstate - prediction.lm)^2))
cat("Root Mean Squared for linear model with all predictors included:",rms.lm)

## Root Mean Squared for linear model with all predictors included: 2133.462

# Calculate test r-squared.
rsquqred <- 1 - error / mean((mytest$Outstate - mean(mytest$Outstate))^2)
cat("Test R squared of 11 predictors GAM model:",rsquqred)

## Test R squared of 11 predictors GAM model: 0.7575958
```

As the output shows, root mean squared (RMS) of 11 predictors GAM model is smaller than the standard linear model with all predictors included. Test R squared is 0.7575958 which means 75.8% of the variability in 'Outstate' in the testing data set is described by the 11 predictors GAM model.